

ug2x7rii4

April 16, 2025

Califonia House Price Dataset:

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv('housing.csv.zip')
df
```

```
[2]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...	
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

	population	households	median_income	median_house_value	\
0	322.0	126.0	8.3252	452600.0	
1	2401.0	1138.0	8.3014	358500.0	
2	496.0	177.0	7.2574	352100.0	
3	558.0	219.0	5.6431	341300.0	
4	565.0	259.0	3.8462	342200.0	
...	
20635	845.0	330.0	1.5603	78100.0	
20636	356.0	114.0	2.5568	77100.0	
20637	1007.0	433.0	1.7000	92300.0	
20638	741.0	349.0	1.8672	84700.0	
20639	1387.0	530.0	2.3886	89400.0	

```

      ocean_proximity
0      NEAR BAY
1      NEAR BAY
2      NEAR BAY
3      NEAR BAY
4      NEAR BAY
...
20635    INLAND
20636    INLAND
20637    INLAND
20638    INLAND
20639    INLAND

```

[20640 rows x 10 columns]

```
[3]: df.head()
```

```

[3]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88           41.0           880.0           129.0
1    -122.22    37.86           21.0          7099.0          1106.0
2    -122.24    37.85           52.0          1467.0           190.0
3    -122.25    37.85           52.0          1274.0           235.0
4    -122.25    37.85           52.0          1627.0           280.0

      population  households  median_income  median_house_value  ocean_proximity
0         322.0        126.0         8.3252         452600.0      NEAR BAY
1        2401.0       1138.0         8.3014        358500.0      NEAR BAY
2         496.0        177.0         7.2574        352100.0      NEAR BAY
3         558.0        219.0         5.6431        341300.0      NEAR BAY
4         565.0        259.0         3.8462        342200.0      NEAR BAY

```

```
[4]: df.tail()
```

```

[4]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
20635    -121.09    39.48           25.0          1665.0           374.0
20636    -121.21    39.49           18.0           697.0           150.0
20637    -121.22    39.43           17.0          2254.0           485.0
20638    -121.32    39.43           18.0          1860.0           409.0
20639    -121.24    39.37           16.0          2785.0           616.0

      population  households  median_income  median_house_value  \
20635         845.0        330.0         1.5603          78100.0
20636         356.0        114.0         2.5568          77100.0
20637        1007.0        433.0         1.7000          92300.0
20638         741.0        349.0         1.8672          84700.0
20639        1387.0        530.0         2.3886          89400.0

```

```

ocean_proximity
20635      INLAND
20636      INLAND
20637      INLAND
20638      INLAND
20639      INLAND

```

```
[5]: df.shape
```

```
[5]: (20640, 10)
```

```
[6]: df.describe()
```

```
[6]:
```

	longitude	latitude	housing_median_age	total_rooms \
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081
std	2.003532	2.135952	12.585558	2181.615252
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639

```

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object

dtypes: float64(9), object(1)

memory usage: 1.6+ MB

```
[8]: df.isnull().sum()
```

```
[8]: longitude          0
latitude             0
housing_median_age   0
total_rooms          0
total_bedrooms      207
population           0
households           0
median_income        0
median_house_value   0
ocean_proximity      0
dtype: int64
```

```
[9]: df['total_bedrooms'].fillna(df['total_bedrooms'].mean(), inplace=True)
df.isnull().sum()
```

```
[9]: longitude          0
latitude             0
housing_median_age   0
total_rooms          0
total_bedrooms       0
population           0
households           0
median_income        0
median_house_value   0
ocean_proximity      0
dtype: int64
```

```
[10]: n = df.select_dtypes(exclude='object')
n
```

```
[10]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...	
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

	population	households	median_income	median_house_value
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0
3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0
...
20635	845.0	330.0	1.5603	78100.0
20636	356.0	114.0	2.5568	77100.0
20637	1007.0	433.0	1.7000	92300.0
20638	741.0	349.0	1.8672	84700.0
20639	1387.0	530.0	2.3886	89400.0

[20640 rows x 9 columns]

```
[11]: c = df.select_dtypes(include='object')
c
```


```
[11]:
```

	ocean_proximity
0	NEAR BAY
1	NEAR BAY
2	NEAR BAY
3	NEAR BAY
4	NEAR BAY
...	...
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND

[20640 rows x 1 columns]


UNIVARIATE ANALYSIS

```
[12]: plt.figure(figsize=(10,5))
sns.histplot(df['median_house_value'], bins=30, kde=True)
plt.show()
```



```
[13]: sns.countplot(df['ocean_proximity'])
plt.show()

df['ocean_proximity'].value_counts()
```



```
[13]: ocean_proximity
<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: count, dtype: int64
```

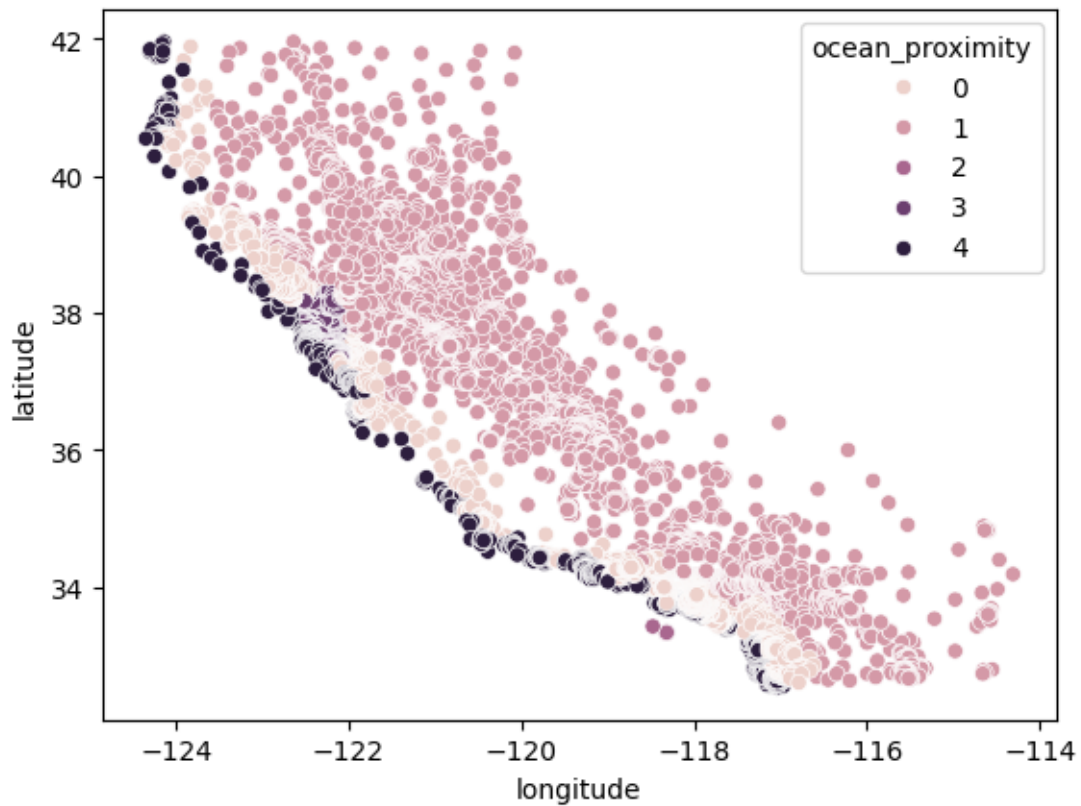
```
[53]: df = df[df['ocean_proximity'] != 'ISLAND']
```

```
[54]: from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
df['ocean_proximity'] = label.fit_transform(df['ocean_proximity'])
df['ocean_proximity']
```

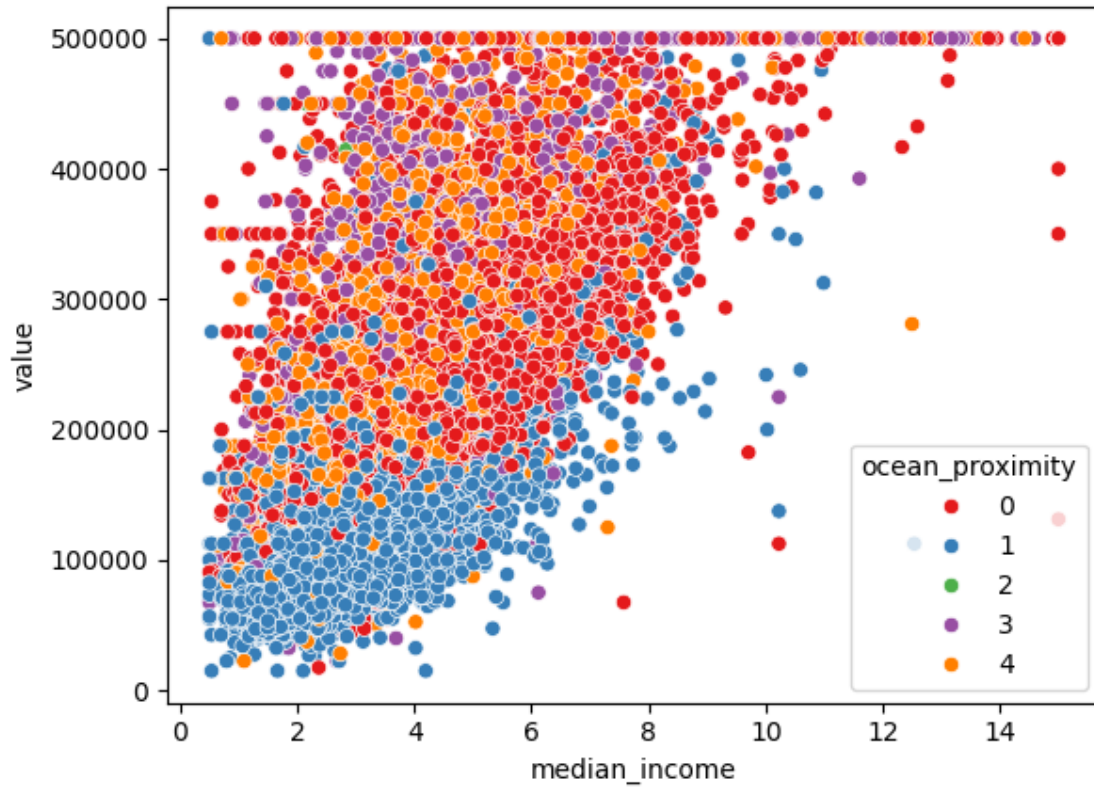
```
[54]: 0      3
1      3
2      3
3      3
4      3
..
20635  1
20636  1
```

```
20637    1
20638    1
20639    1
Name: ocean_proximity, Length: 20640, dtype: int64
```


```
[55]: sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=df['ocean_proximity'])
plt.show()
```



```
[58]: sns.scatterplot(x=df['median_income'], y= df['value'],
    ↪hue=df['ocean_proximity'], palette='Set1')
plt.show()
```

```
[59]: sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=df['value'], palette=
      ↪= 'viridis')
plt.show()
```




```
[19]: c.columns
```

```
[19]: Index(['ocean_proximity'], dtype='object')
```


```
[20]: n.columns
```

```
[20]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
          'total_bedrooms', 'population', 'households', 'median_income',
          'median_house_value'],
          dtype='object')
```

```
[21]: sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=df['households'],
    ↪ palette = 'viridis')
plt.show()
```



```
[22]: sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=df['population'].
      ↪mean(), palette = 'Set2')
plt.show()
```




```
[64]: house_value1 = n.groupby('total_rooms')['median_house_value'].mean().
      ↪reset_index() # Use the actual column name 'median_house_value' instead of
      ↪the undefined variable 'value'.
      house_value1
```

```
[64]:
```

	total_rooms	median_house_value
0	2.0	137500.0
1	6.0	55000.0
2	8.0	500001.0
3	11.0	162500.0
4	12.0	67500.0
...
5921	30450.0	174300.0
5922	32054.0	253900.0
5923	32627.0	118800.0
5924	37937.0	366300.0
5925	39320.0	153700.0

```
[5926 rows x 2 columns]
```

```
[24]: sns.pairplot(house_value1, diag_kind='kde')
plt.show()
```



```
[65]: house_value2 = n.groupby('total_bedrooms')['median_house_value'].mean().
      ↪reset_index()
house_value2
```


```
[65]:
```

	total_bedrooms	median_house_value
0	1.0	500001.000000
1	2.0	96250.000000
2	3.0	285000.000000
3	4.0	186428.714286
4	5.0	238550.000000
...
1919	5290.0	253900.000000
1920	5419.0	500001.000000
1921	5471.0	366300.000000

```
1922      6210.0      153700.000000
1923      6445.0      118800.000000
```

```
[1924 rows x 2 columns]
```


```
[26]: sns.pairplot(house_value2, diag_kind = 'kde', palette = 'viridis')
plt.show()
```




```
[66]: house_value3 = n.groupby('population')['median_house_value'].mean().
      ↪reset_index()
house_value3
population = df['population'].sample(n=50)
house_value3.plot(x='population', y='median_house_value', kind='bar',
      ↪color='green')
plt.xlabel('Population')
plt.ylabel('Median House Value')
plt.title('Population vs Median House Value')
```

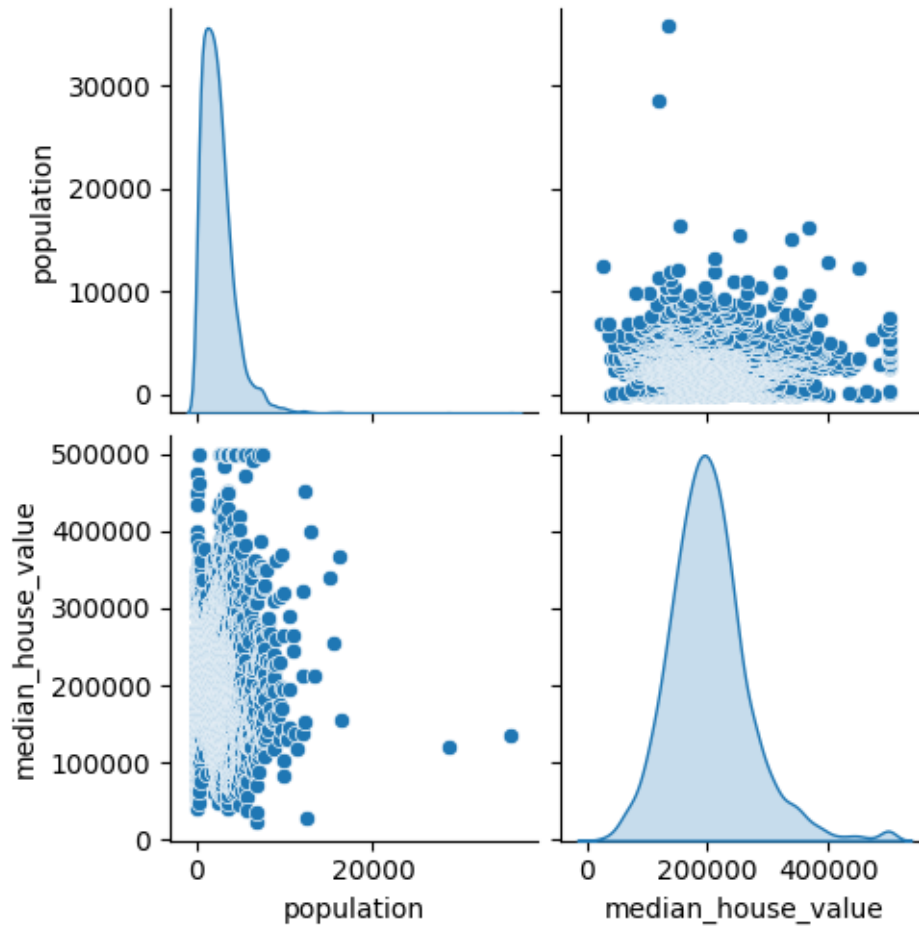
```
plt.xticks([])
plt.show()

sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=df['population'],
               palette = 'viridis')
plt.show()
```





```
[28]: sns.pairplot(house_value3, diag_kind = 'kde', palette = 'viridis')  
plt.show()
```





```
[29]: n.columns
```

```
[29]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
          'total_bedrooms', 'population', 'households', 'median_income',
          'median_house_value'],
          dtype='object')
```

```
[30]: pop_income = n.groupby('population')['median_income'].mean().reset_index()
      pop_income

      pop_income = pop_income.sample(n=50)
      pop_income.plot(x='population', y='median_income', kind='bar')
      plt.xticks([])
      plt.show()
```



```
[31]: pop_households = n.groupby('population')['households'].mean().reset_index()
pop_households

pop_households = pop_households.sample(n=50)
pop_households.plot(x='population', y='households', kind='bar')
plt.xticks([])
plt.show()


# remove outliers


q1 = pop_households['households'].quantile(0.25)
q3 = pop_households['households'].quantile(0.75)
iqr = q3 - q1
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr
pop_households = pop_households[(pop_households['households'] >= lower_bound) &
    ↪ (pop_households['households'] <= upper_bound)]
pop_households


outliers = pop_households[pop_households['households'] > lower_bound ]
outliers.plot(x='population', y='households', kind='bar')
plt.xlabel('Population')
```

```
plt.ylabel('Households')
plt.title('Outliers Population vs Households')
plt.xticks([])
plt.show()


pop_households.plot(x='population', y='households', kind='bar')
plt.xlabel('Population')
plt.ylabel('Households')
plt.title('Population vs Households')
plt.xticks([])
plt.show()
```







```
[32]: sns.pairplot(pop_households, diag_kind = 'kde', palette = 'viridis')  
plt.show()
```



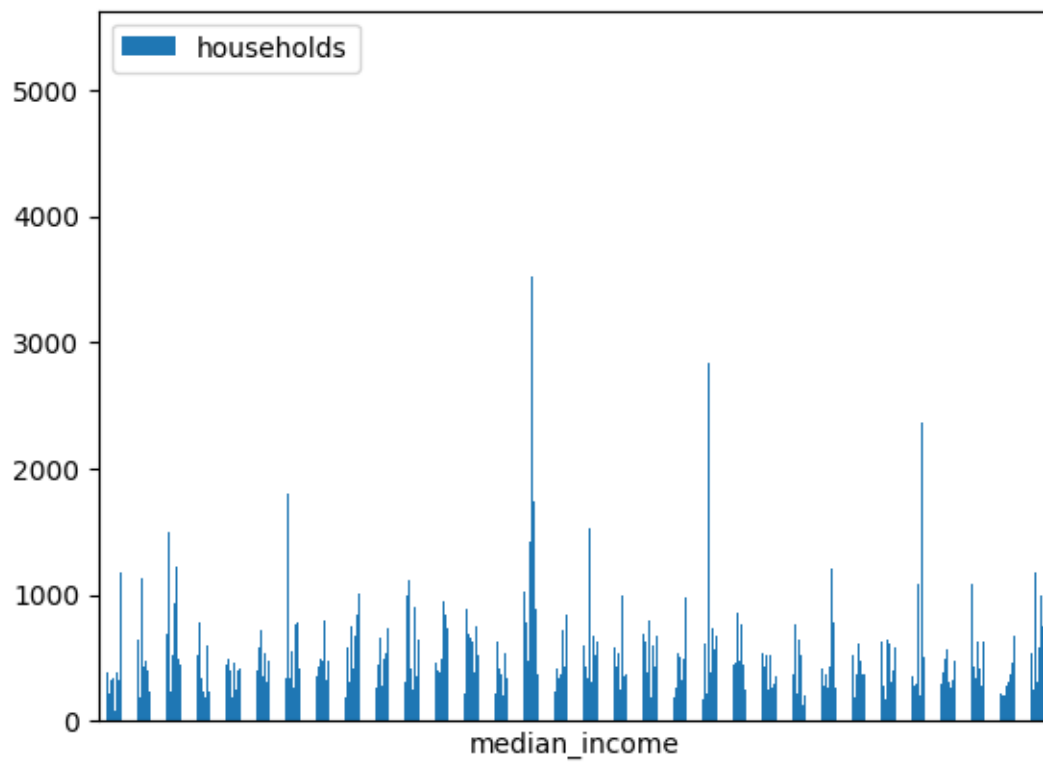
```
[33]: n.columns
```


```
[33]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
            'total_bedrooms', 'population', 'households', 'median_income',
            'median_house_value'],
            dtype='object')
```

```
[34]: income_households = n.groupby('median_income')['households'].mean().
      ↪reset_index()
income_households.plot(x='median_income', y='households', kind='bar')
plt.xticks([])
plt.show()


income_households = income_households.sample(n=50)
sns.barplot(x='median_income', y='households', data=income_households)
plt.xlabel('Median Income')
plt.ylabel('Households')
```

```
plt.title('Median Income vs Households')  
plt.xticks([])  
plt.show()
```





```
[35]: corr_matrix = n.corr()  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')  
plt.show()
```

```
[36]: corr_matrix = corr_matrix.style.background_gradient(cmap='coolwarm')
corr_matrix
```

```
[36]: <pandas.io.formats.style.Styler at 0x7fe9dc6cf7d0>
```

```
[37]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import *
from sklearn.metrics import *
```

```
[38]: from sklearn import *
```

```
[39]: df = df.rename(columns={'median_house_value': 'value'})
df
```

```
[39]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	
...	
20635	-121.09	39.48	25.0	1665.0	374.0	
20636	-121.21	39.49	18.0	697.0	150.0	
20637	-121.22	39.43	17.0	2254.0	485.0	
20638	-121.32	39.43	18.0	1860.0	409.0	
20639	-121.24	39.37	16.0	2785.0	616.0	

	population	households	median_income	value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	3
1	2401.0	1138.0	8.3014	358500.0	3
2	496.0	177.0	7.2574	352100.0	3
3	558.0	219.0	5.6431	341300.0	3
4	565.0	259.0	3.8462	342200.0	3
...
20635	845.0	330.0	1.5603	78100.0	1
20636	356.0	114.0	2.5568	77100.0	1
20637	1007.0	433.0	1.7000	92300.0	1
20638	741.0	349.0	1.8672	84700.0	1
20639	1387.0	530.0	2.3886	89400.0	1

[20640 rows x 10 columns]

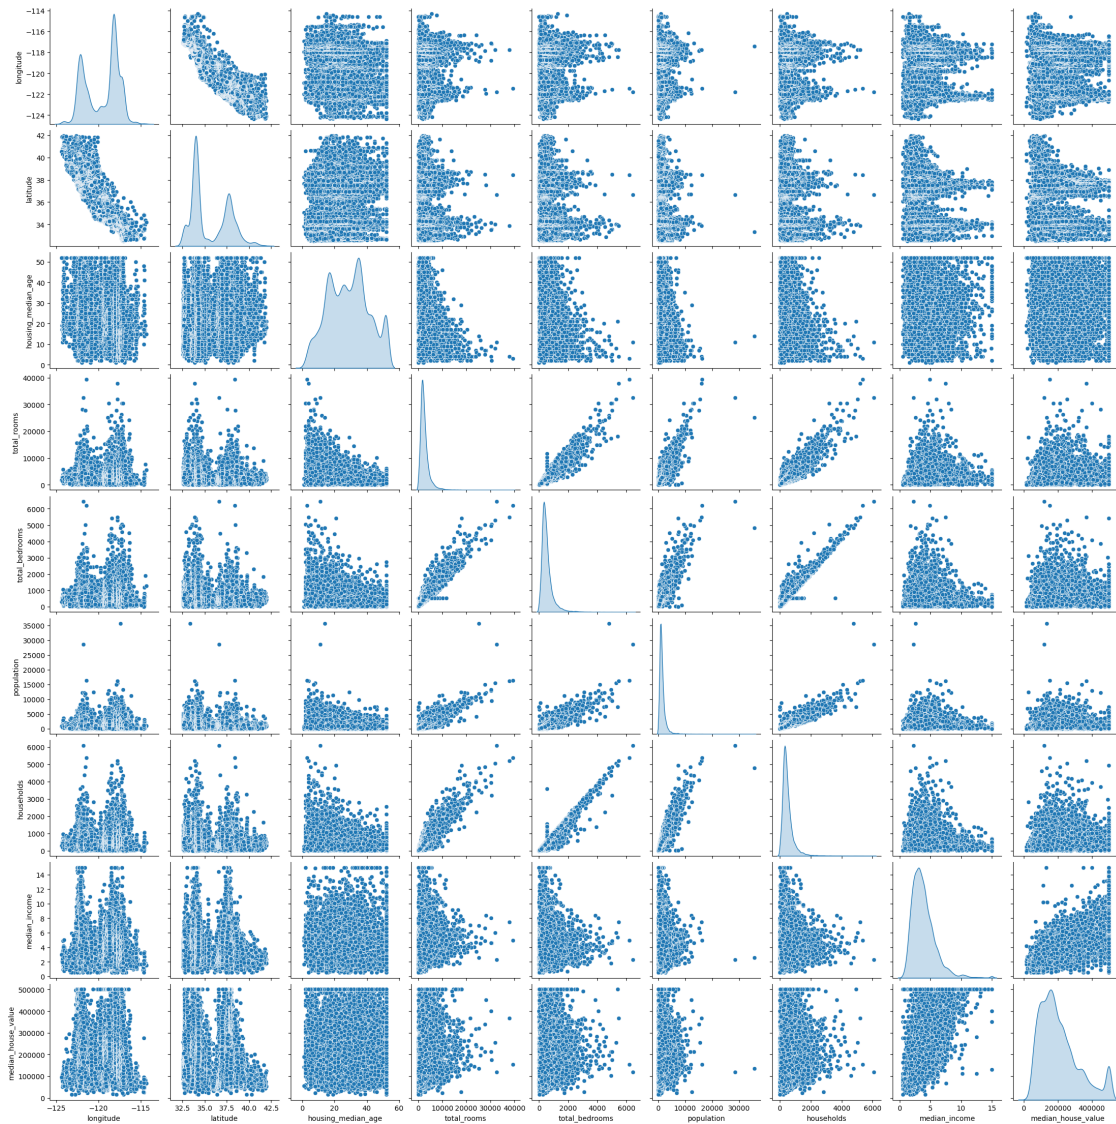
```
[40]: label_encoder = LabelEncoder()
df['ocean_proximity'] = label_encoder.fit_transform(df['ocean_proximity'])
op = df['ocean_proximity']
op
```

```
[40]:
```


0	3
1	3
2	3
3	3
4	3
...	
20635	1
20636	1
20637	1
20638	1
20639	1

Name: ocean_proximity, Length: 20640, dtype: int64

```
[41]: sns.pairplot(n, diag_kind='kde')  
plt.show()
```



```
[42]: sns.scatterplot(x=df['longitude'], y= df['latitude'], hue=op, palette =  
    ↪ 'viridis')  
plt.show()
```



```
[43]: x = df.drop('value', axis=1)
      y = df['value']

      x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      random_state=42)

      model = LinearRegression()
      model.fit(x_train, y_train)


      y_pred = model.predict(x_test)

      mse = mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)

      print("Mean Squared Error:", mse)
      print("R-squared:", r2)
```

Mean Squared Error: 5055025116.165619
R-squared: 0.6142406531011781

```
[44]: sns.lineplot(x=y_test, y=y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted', color = 'green')
plt.show()
```



```
[45]: from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score

model = DecisionTreeRegressor()
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

Mean Squared Error: 4634878555.809109

R-squared: 0.6463028998755034

```
[46]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report
      from sklearn.preprocessing import StandardScaler
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import *
```

```
[47]: x = df.drop('ocean_proximity', axis=1)
      y = df['ocean_proximity']

      x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42)

      scaler = StandardScaler()
      x_train_scaled = scaler.fit_transform(x_train)
      x_test_scaled = scaler.transform(x_test)

      model = LogisticRegression()
      model.fit(x_train_scaled, y_train)

      y_pred = model.predict(x_test_scaled)

      accuracy = accuracy_score(y_test, y_pred)
      classification_rep = classification_report(y_test, y_pred)

      print("Accuracy:", accuracy)
      print("Classification Report:\n", classification_rep)
```


Accuracy: 0.7950581395348837

Classification Report:

	precision	recall	f1-score	support
0	0.74	0.92	0.82	1795
1	0.97	0.91	0.94	1324
2	0.00	0.00	0.00	1
3	0.65	0.76	0.70	436
4	0.63	0.17	0.26	572
accuracy			0.80	4128
macro avg	0.60	0.55	0.55	4128
weighted avg	0.79	0.80	0.77	4128

```
[48]: sns.lineplot(x=y_test, y=y_pred)
      plt.xlabel('Actual Values')
      plt.ylabel('Predicted')
```

```
plt.title('Actual vs Predicted', color = 'green')
plt.show()
```



```
[49]: x = df.drop('total_rooms', axis=1)
      y = df['total_rooms']

      x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
        ↪random_state=42)

      scaler = StandardScaler()
      x_train_scaled = scaler.fit_transform(x_train)
      x_test_scaled = scaler.transform(x_test)

      model = LinearRegression()
      model.fit(x_train_scaled, y_train)

      y_pred = model.predict(x_test_scaled)

      mse = mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
print("R-squared:", r2)

sns.lineplot(x=y_test, y=y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted', color = 'green')
plt.show()
```

Mean Squared Error: 494760.2558706711

R-squared: 0.8986164708974816

